

NAME: AVIREDDY NVSRK ROHAN

REG.NO: 19BCE1180

DATASET USED: KDD CUP

```
In [1]: import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from matplotlib import pyplot as plt
```

```
In [2]: from keras.layers import Dense
from keras.models import Sequential
```

```
In [3]: data=pd.read_csv("/content/drive/MyDrive/ML_LAB/corrected",header=None)
data.columns=['duration','protocol_type','service','flag','src_bytes','dst_bytes','land
```

```
In [4]: data.head()
```

```
Out[4]:
```

| | duration | protocol_type | service | flag | src_bytes | dst_bytes | land | wrong_fragment | urgent | hot | nu |
|---|----------|---------------|---------|------|-----------|-----------|------|----------------|--------|-----|----|
| 0 | 0 | udp | private | SF | 105 | 146 | 0 | 0 | 0 | 0 | |
| 1 | 0 | udp | private | SF | 105 | 146 | 0 | 0 | 0 | 0 | |
| 2 | 0 | udp | private | SF | 105 | 146 | 0 | 0 | 0 | 0 | |
| 3 | 0 | udp | private | SF | 105 | 146 | 0 | 0 | 0 | 0 | |
| 4 | 0 | udp | private | SF | 105 | 146 | 0 | 0 | 0 | 0 | |

```
In [ ]: data["class"].value_counts()
```

```
Out[ ]:
```

| | |
|----------------|--------|
| smurf. | 164091 |
| normal. | 60593 |
| neptune. | 58001 |
| snmpgetattack. | 7741 |
| mailbomb. | 5000 |
| guess_passwd. | 4367 |
| snmpguess. | 2406 |
| satan. | 1633 |
| warezmaster. | 1602 |
| back. | 1098 |
| mscan. | 1053 |
| apache2. | 794 |

```

processtable.      759
saint.             736
portsweep.         354
ipsweep.           306
httptunnel.        158
pod.               87
nmap.              84
buffer_overflow.   22
multihop.          18
sendmail.          17
named.             17
ps.               16
rootkit.           13
xterm.             13
teardrop.          12
land.              9
xlock.             9
xsnoop.            4
ftp_write.         3
udpstorm.          2
perl.              2
sqlattack.         2
loadmodule.        2
worm.              2
phf.               2
imap.              1
Name: class, dtype: int64

```

```

In [5]: le=LabelEncoder()
data["protocol_type"]=le.fit_transform(data["protocol_type"])
data["service"]=le.fit_transform(data["service"])
data["flag"]=le.fit_transform(data["flag"])
data["protocol_type"]=data["protocol_type"].astype("object")
data["service"]=data["service"].astype("object")
data["flag"]=data["flag"].astype("object")
data["is_host_login"]=data["is_host_login"].astype("object")
data["is_guest_login"]=data["is_guest_login"].astype("object")
data["logged_in"]=data["logged_in"].astype("object")

```

```

In [6]: cond=((data['class']=="smurf.") | (data['class']=="normal.") | (data['class']=="satan."

```

```

In [7]: new_data=data[cond]

```

```

In [8]: multi_label=new_data["class"]
new_data.drop(["class"],inplace=True,axis=1)
multi_label=le.fit_transform(multi_label)

```

/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py:4174: SettingWithCopyWarning:
g:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
errors=errors,

```

In [ ]: new_data.columns.get_loc("root_shell")

```

Out[]: 13

```
In [9]: new_data.shape
```

Out[9]: (227053, 41)

```
In [ ]: temp=pd.DataFrame()
for i in new_data.columns:
    if new_data.dtypes[i] is not 'O':
        temp.insert(new_data.columns.get_loc(i),i,new_data[i])
```

```
In [10]: st_obj=StandardScaler()
data_std=st_obj.fit_transform(new_data)
data_std=pd.DataFrame(data_std)
```

```
In [11]: import imblearn
from imblearn.over_sampling import SMOTE
```

/usr/local/lib/python3.7/dist-packages/sklearn/externals/six.py:31: FutureWarning: The module is deprecated in version 0.21 and will be removed in version 0.23 since we've dropped support for Python 2.7. Please rely on the official version of six (<https://pypi.org/project/six/>).

"(<https://pypi.org/project/six/>).", FutureWarning)

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:144: FutureWarning: The sklearn.neighbors.base module is deprecated in version 0.22 and will be removed in version 0.24. The corresponding classes / functions should instead be imported from sklearn.neighbors. Anything that cannot be imported from sklearn.neighbors is now part of the private API.

warnings.warn(message, FutureWarning)

class distributions before using SMOTE for 80:20 train test split

```
In [13]: train_data,test_data,train_labels,test_labels=train_test_split(data_std,multi_label,tes
```

```
In [ ]: g=pd.Series(train_labels)
```

```
In [ ]: g.value_counts()*100/train_labels.shape[0]
```

```
Out[ ]: 3    72.292752
0     26.671144
2      0.718997
1      0.317107
dtype: float64
```

Multilayer perceptron model

```
In [33]: mlp=Sequential()
mlp.add(Dense(32,input_dim=41,activation="relu"))
```

```
mlp.add(Dense(16,activation="relu"))
mlp.add(Dense(8,activation="relu"))
mlp.add(Dense(1,activation="sigmoid"))
mlp.compile(optimizer="Adam",loss="categorical_crossentropy",metrics=['accuracy'])
```

Performance before SMOTE

```
In [44]: mlp.fit(train_data,train_labels)
```

```
5677/5677 [=====] - 12s 2ms/step - loss: 0.0000e+00 - accuracy: 0.2606
```

```
Out[44]: <keras.callbacks.History at 0x7f29f898b810>
```

on training data

```
In [45]: uns_train_pred=mlp.predict(train_data)
```


```
In [46]: print(classification_report(train_labels,uns_train_pred))
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1272: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.97 | 0.99 | 48446 |
| 1 | 0.00 | 0.86 | 0.01 | 576 |
| 2 | 0.00 | 0.00 | 0.00 | 1306 |
| 3 | 0.00 | 0.00 | 0.00 | 131314 |
| accuracy | | | 0.26 | 181642 |
| macro avg | 0.25 | 0.46 | 0.25 | 181642 |
| weighted avg | 0.27 | 0.26 | 0.26 | 181642 |

```
In [54]: cm=confusion_matrix(train_labels,uns_train_pred)
disp=ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=["0","1","2","3"])
disp.plot()
plt.show()
```



on testing data


```
In [48]: uns_test_pred=mlp.predict(test_data)
```

```
In [49]: print(classification_report(test_labels,uns_test_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.97 | 0.98 | 12147 |
| 1 | 0.00 | 0.82 | 0.01 | 160 |
| 2 | 0.00 | 0.00 | 0.00 | 327 |
| 3 | 0.00 | 0.00 | 0.00 | 32777 |
| accuracy | | | 0.26 | 45411 |
| macro avg | 0.25 | 0.45 | 0.25 | 45411 |
| weighted avg | 0.27 | 0.26 | 0.26 | 45411 |

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1272: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

```
In [55]: cm=confusion_matrix(test_labels,uns_test_pred)
disp=ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=["0","1","2","3"])
disp.plot()
plt.show()
```



PERFORMANCE AFTER SMOTE

```
In [12]: smote=SMOTE()
X_smote,y_smote=smote.fit_resample(data_std,multi_label)
```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be removed in version 0.24.

warnings.warn(msg, category=FutureWarning)

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be removed in version 0.24.

warnings.warn(msg, category=FutureWarning)

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be removed in version 0.24.

warnings.warn(msg, category=FutureWarning)

```
In [50]: train_smote,test_smote,train_smote_labs,test_smote_labs=train_test_split(X_smote,y_smote,
```

```
In [51]: mlp.fit(train_smote,train_smote_labs)
```

18461/18461 [=====] - 28s 2ms/step - loss: 0.0000e+00 - accuracy: 0.4544

```
Out[51]: <keras.callbacks.History at 0x7f29f0833690>
```

on smote train data

```
In [52]: mlp_train_pred=mlp.predict(train_smote)
```


```
In [53]: print(classification_report(train_smote_labs,mlp_train_pred))
```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1272: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.87 | 0.96 | 0.91 | 147712 |
| 1 | 0.30 | 0.86 | 0.44 | 147373 |
| 2 | 0.00 | 0.00 | 0.00 | 147757 |
| 3 | 0.00 | 0.00 | 0.00 | 147885 |
| accuracy | | | 0.45 | 590727 |
| macro avg | 0.29 | 0.45 | 0.34 | 590727 |
| weighted avg | 0.29 | 0.45 | 0.34 | 590727 |

```
In [56]: cm=confusion_matrix(train_smote_labs,mlp_train_pred)
disp=ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=["0","1","2","3"])
disp.plot()
plt.show()
```



on smote test data

```
In [58]: mlp_test_pred=mlp.predict(test_smote)
```

```
In [59]: print(classification_report(test_smote_labs,mlp_test_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.87 | 0.96 | 0.91 | 16379 |
| 1 | 0.30 | 0.86 | 0.45 | 16718 |
| 2 | 0.00 | 0.00 | 0.00 | 16334 |
| 3 | 0.00 | 0.00 | 0.00 | 16206 |
| accuracy | | | 0.46 | 65637 |
| macro avg | 0.29 | 0.45 | 0.34 | 65637 |
| weighted avg | 0.29 | 0.46 | 0.34 | 65637 |

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1272: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
```

```
In [60]: cm=confusion_matrix(test_smote_labs,mlp_test_pred)
```

```
disp=ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=["0","1","2","3"])  
disp.plot()  
plt.show()
```

