

AVIREDDY NVSRK ROHAN

19BCE1180

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
from matplotlib import cm
from sklearn.metrics import silhouette_samples
```

```
In [2]: iris_data = pd.read_csv("/content/iris.csv")
X = iris_data[[column for column in iris_data.columns if column != 'species']]
y = iris_data['species']
```

```
In [3]: iris_data.columns
```

```
Out[3]: Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
              'species'],
              dtype='object')
```

clustering by kmeans++

```
In [ ]: iris_data
```

```
Out[10]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa
...
145	6.7	3.0	5.2	2.3	virginica
146	6.3	2.5	5.0	1.9	virginica
147	6.5	3.0	5.2	2.0	virginica
148	6.2	3.4	5.4	2.3	virginica
149	5.9	3.0	5.1	1.8	virginica


150 rows × 5 columns

```

In [4]: km = KMeans(n_clusters=3,
                    init='k-means++',
                    n_init=10,
                    max_iter=300,
                    tol=1e-04,
                    random_state=0)
y_km = km.fit_predict(X)
cluster_labels = np.unique(y_km)
n_clusters = cluster_labels.shape[0]
silhouette_vals = silhouette_samples(X, y_km, metric='euclidean')
y_ax_lower, y_ax_upper = 0, 0
yticks = []
for i, c in enumerate(cluster_labels):
    c_silhouette_vals = silhouette_vals[y_km == c]
    c_silhouette_vals.sort()
    y_ax_upper += len(c_silhouette_vals)
    color = cm.jet(float(i) / n_clusters)
    plt.barh(range(y_ax_lower, y_ax_upper), c_silhouette_vals, height=1.0,
             edgecolor='none', color=color)
    yticks.append((y_ax_lower + y_ax_upper) / 2.)
    y_ax_lower += len(c_silhouette_vals)

silhouette_avg = np.mean(silhouette_vals)
plt.axvline(silhouette_avg, color="red", linestyle="--")
plt.yticks(yticks, cluster_labels + 1)
plt.ylabel('Cluster')
plt.xlabel('Silhouette coefficient')
plt.tight_layout()
# plt.savefig('./figures/silhouette.png', dpi=300)
plt.show()


```



```

In [5]: km = KMeans(n_clusters=3,
                    init='random',
                    n_init=1,
                    max_iter=20,
                    tol=1e-04,
                    random_state=0)
y_km = km.fit_predict(X)
X = np.array(X)
plt.scatter(X[y_km == 0,0],
            X[y_km == 0,1],
            s=50,
            c='blue',
            marker='s',
            label='cluster 1')
plt.scatter(X[y_km == 1,0],
            X[y_km == 1,1],
            s=50,
            c='orange',
            marker='o',
            label='cluster 2')
plt.scatter(X[y_km == 2,0],
            X[y_km == 2,1],
            s=50,
            c='yellow',
            marker='o',
            label='cluster 3')
plt.scatter(km.cluster_centers_[0, 0],
            km.cluster_centers_[0, 1],
            s=250,
            marker='*',
            c='red',
            label='centroids')
plt.legend()
plt.grid()
plt.tight_layout()
#plt.savefig('./figures/centroids.png', dpi=300)
plt.show()

```



cluster labels

```
In [ ]: y_km
```

```
Out[12]: array([2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
                2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
                2, 2, 2, 2, 2, 2, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
                1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
                1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
                0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0,
                0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1], dtype=int32)
```

CLUSTER VS ACTUAL LABELS

```
In [9]: y_num = [2 if label == 'Iris-setosa' else 1 if label == 'Iris-versicolor' else 0
```

```
In [11]: from sklearn.metrics import classification_report
cr = classification_report(y_num, y_km, digits=3)
print("Classification Report\n\n", cr)
```

Classification Report

	precision	recall	f1-score	support
0	1.000	0.260	0.413	150
1	0.000	0.000	0.000	0
2	0.000	0.000	0.000	0
accuracy			0.260	150
macro avg	0.333	0.087	0.138	150
weighted avg	1.000	0.260	0.413	150

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

model buliding

```
In [6]: from sklearn.linear_model import LogisticRegression
```

```
In [12]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y_km, test_size=0.3, random_state=42)

from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
sc.fit(X_train)
X_train_std = sc.transform(X_train)
X_test_std = sc.transform(X_test)
```

```
In [13]: model = LogisticRegression()
model.fit(X_train_std, y_train)

y_pred = model.predict(X_test_std)
print('Total Examples: {}\nMisclassified examples: {}'.format((y_test.size), (y_t

Total Examples: 45
Misclassified examples: 1
```

```
In [16]: cr = classification_report(y_test, y_pred, digits=3)
print(cr)
```


	precision	recall	f1-score	support
0	1.000	0.909	0.952	11
1	0.950	1.000	0.974	19
2	1.000	1.000	1.000	15
accuracy			0.978	45
macro avg	0.983	0.970	0.976	45
weighted avg	0.979	0.978	0.978	45

```
In [17]: cm = skm.plot_confusion_matrix(model, X_test_std, y_test, cmap='Blues', values_fc
print(cm)
```

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7fa2e2a51f50>

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated; Function `plot_confusion_matrix` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: ConfusionMatrixDisplay.from_predictions or ConfusionMatrixDisplay.from_estimator.


warnings.warn(msg, category=FutureWarning)



Gaussian mixture clustering

```
In [8]: from sklearn.mixture import GaussianMixture
```

```
In [18]: sc = GaussianMixture(n_components=3, covariance_type='full')
preds = sc.fit_predict(X)
X = np.array(X)
plt.scatter(X[preds == 0,0],
            X[preds == 0,1],
            s=50,
            c='blue',
            marker='s',
            label='cluster 1')
plt.scatter(X[preds == 1,0],
            X[preds == 1,1],
            s=50,
            c='orange',
            marker='o',
            label='cluster 2')
plt.scatter(X[preds == 2,0],
            X[preds == 2,1],
            s=50,
            c='yellow',
            marker='o',
            label='cluster 3')
plt.scatter(km.cluster_centers[:, 0],
            km.cluster_centers[:, 1],
            s=250,
            marker='*',
            c='red',
            label='centroids')
plt.legend()
plt.grid()
plt.tight_layout()
#plt.savefig('./figures/centroids.png', dpi=300)
plt.show()
```



cluster labels to classifier

```
print('Cluster labels: %s' % preds)
```

[illegible]

```
y_num2 = [0 if label == 'Iris-setosa' else 1 if label == 'Iris-versicolor' else 2 if label == 'Iris-virginica' else 0
print(y_num2[:10])
print(preds[:10])
```

$$\begin{bmatrix} 2 & 2 & 2 & 2 & 2 & 2 & 2 & 2 & 2 & 2 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

cluster predicted vs actual labels

```
import sklearn.metrics as skm
cr = skm.classification_report(y_num2, preds, digits=3)
print("Classification Report\n\n", cr)
```

Classification Report

	precision	recall	f1-score	support
0	0.000	0.000	0.000	0
1	0.000	0.000	0.000	0
2	1.000	0.367	0.537	150
accuracy			0.367	150
macro avg	0.333	0.122	0.179	150
weighted avg	1.000	0.367	0.537	150

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308:
UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0
in labels with no true samples. Use `zero_division` parameter to control this b
ehavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308:
UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0
in labels with no true samples. Use `zero_division` parameter to control this b
ehavior.
```

```
warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308:
UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0
in labels with no true samples. Use `zero_division` parameter to control this b
ehavior.
```

```
warn prf(average, modifier, msg start, len(result))
```



```
In [24]: X_train, X_test, y_train, y_test = train_test_split(X, preds, test_size=0.3, rand
```

```
sc = StandardScaler()
sc.fit(X_train)
X_train_std = sc.transform(X_train)
X_test_std = sc.transform(X_test)

model2 = LogisticRegression()
model2.fit(X_train_std, y_train)

y_pred = model2.predict(X_test_std)
print('Total Examples: {}\nMisclassified examples: {}'.format((y_test.size), (y_t
```

```
Total Examples: 45
Misclassified examples: 1
```

```
In [25]: cr = classification_report(y_test, y_pred, digits=3)
print(cr)
```

	precision	recall	f1-score	support
0	1.000	1.000	1.000	15
1	0.933	1.000	0.966	14
2	1.000	0.938	0.968	16
accuracy			0.978	45
macro avg	0.978	0.979	0.978	45
weighted avg	0.979	0.978	0.978	45

```
In [28]: cm = skm.plot_confusion_matrix(model2, X_test_std, y_test, cmap='Blues', values_format='%.1f')
print(cm)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7fa2e07e4090>
```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated; Function `plot_confusion_matrix` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: ConfusionMatrixDisplay.from_predictions or ConfusionMatrixDisplay.from_estimator.

```
warnings.warn(msg, category=FutureWarning)
```

